

Applied Machine Learning

Course number: W207

Prof. Alexander I. Iliev, Ph.D.

Applied Machine Learning

Lecture 8 ...

- *Project requirements*
- *SVM (whiteboard example)*
- *Comparison of ML algorithms discussed in class*
- *Fourier transform, DFT, FFT and IFFT, variations*
- *Audio, Speech, Phonemes, Formants, etc.*

Applied Machine Learning

- *Groups ready by next class on Monday – L9*
- *Project grading criteria:*

Late	Sensible methods (25)	Feature Engineer ing (20)	Error Analysis (15)	Notebook Presentat ion (10)	Problem Descripti on (10)	Good results (10)	In Class Presenta tion (10)
------	-----------------------------	---------------------------------	---------------------------	-----------------------------------	---------------------------------	-------------------------	-----------------------------------

Scipy Modules

- Here is a list of some of the most common Scipy modules:
 - `Scipy.fftpack` – contains Fast Fourier Transforms (FFTs)
 - `Scipy.integrate` – contains a variety of **integrating functions**
 - `Scipy.interpolate` – contains a variety of **interpolation classes**
 - `Scipy.io` – functions for **reading and writing data** from/to a variety of file formats
 - `Scipy.io.wavfile` – read/write data from/to a variety of file formats '**wav**', '**arff**', etc.
 - `Scipy.linalg` – contains linear algebra routines
 - `Scipy.ndimage` – contains many functions for multi-dimensional **image processing**
 - `Scipy.signal` – rich **filtering** capabilities, **wavlets**, **spectral analysis**, and much more
 - `Scipy.optimize` – local **optimization package** and root finding
 - `Scipy.spatial` – **nearest neighbor** queries and **distance functions**
 - `Scipy.stats` – large number of **probability distributions** and **statistical functions**
 - `Scipy.special` – large variety of functions such as: **elliptic**, **bessel**, **legendre**, etc.
 - `Scipy.misc` – variety of other functions

Scipy FFT

- The Scipy FFTpack
 - `scipy.fftpack` vs `numpy.fft`:
 - the `scipy.fftpack` does much more on top of what `numpy.fft` offers:
 - » `fft` and `ifft` - the Discrete Fourier Transform and its inverse of real or complex sequence of numbers
 - » `fft2` and `ifft2` - 2D discrete Fourier transform and its inverse
 - » `fftn` and `ifftn` - multidimensional discrete Fourier transform and its inverse
 - » `dct` and `idct` - Discrete Cosine Transform of arbitrary type sequence
 - » `dst` and `idst` - Discrete Sine Transform of arbitrary type sequence
 - » `tilbert` and `itilbert` - the h-Tilbert transform of a periodic sequence and its inverse
 - » `hilbert` and `ihilbert` - Hilbert Transform of a periodic sequence and its inverse
 - » `fftfreq` - the Discrete Fourier Transform sample frequencies
 - » `convolve` - performs convolution on a given signal
 - » ... and more

Working with files

- Working with files – sound 1/4

reading .wav files

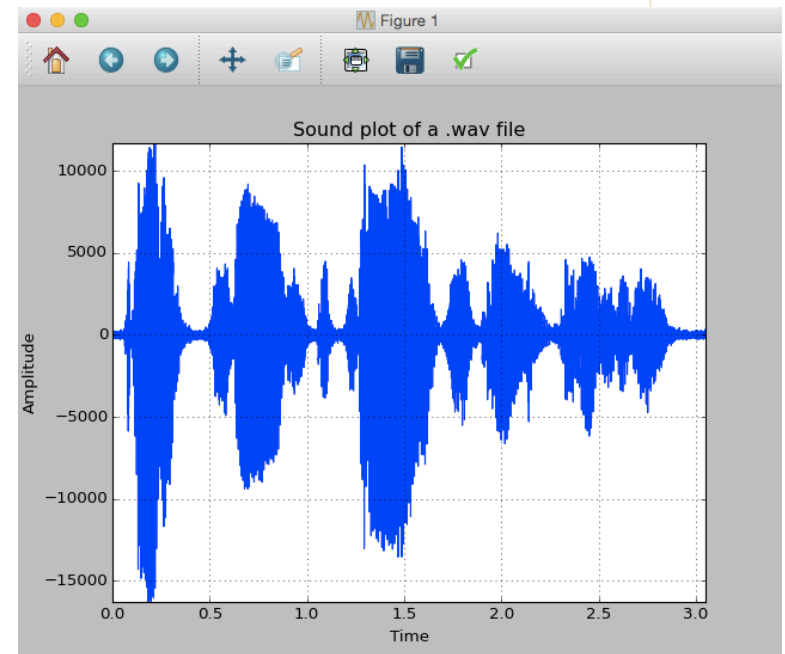
```
Python
In [35]: from scipy.io.wavfile import read
In [36]: (fsx, x) = read('files/lecture8/alex_mono.wav')
In [37]: print(len(x.shape))    # '1' is mono
1
In [38]: print(x[:,])    # to access the channel no digit is used after ','
[-4 43 23 ..., 27 42 -3]
In [39]: x
Out[39]: array([-4, 43, 23, ..., 27, 42, -3], dtype=int16)
In [40]: (fsy, y) = read('files/lecture8/alex_stereo.wav')
In [41]: print(len(y.shape))    # '2' is stereo 2-dimensional array
2
In [42]: y
Out[42]:
array([[ -4,  -4],
       [44, 43],
       [20, 23],
       ...,
       [26, 29],
       [43, 41],
       [-4, -3]], dtype=int16)
In [43]: print(y[:,0])    # to access each channel separately use '0' or '1'
[-4 44 20 ..., 26 43 -4]
In [44]: print(y[:,1])
[-4 43 23 ..., 29 41 -3]
```

Working with files

- Working with files – sound 2/4

plotting .wav files

```
79 ## More on sound:
80 # Example 1 - plotting sounds:
81 from pylab import linspace, plot, title, xlabel, ylabel, grid, axis, pause
82 from scipy.io.wavfile import read
83
84 (Fs, x) = read('files/lecture8/alex_mono.wav') # Fs - sampling frequency, x - signal
85 length = len(x) # number of samples in 'x'
86 time = length/Fs # calculate the length of the .wav file in secs
87 t = linspace(0,time,length) # create evenly spaced numbers between [0:time]
88 plot(t,x) # plot signal 'x'
89 title('Sound plot of a .wav file')
90 xlabel('Time')
91 ylabel('Amplitude')
92 axis('tight')
93 grid(True)
```



Working with files

- Working with files – sound 3/4
 - lets create a (pseudo) stereo music from a single (mono) channel

manipulating
.wav files

Note: this is not
a true stereo
Signal

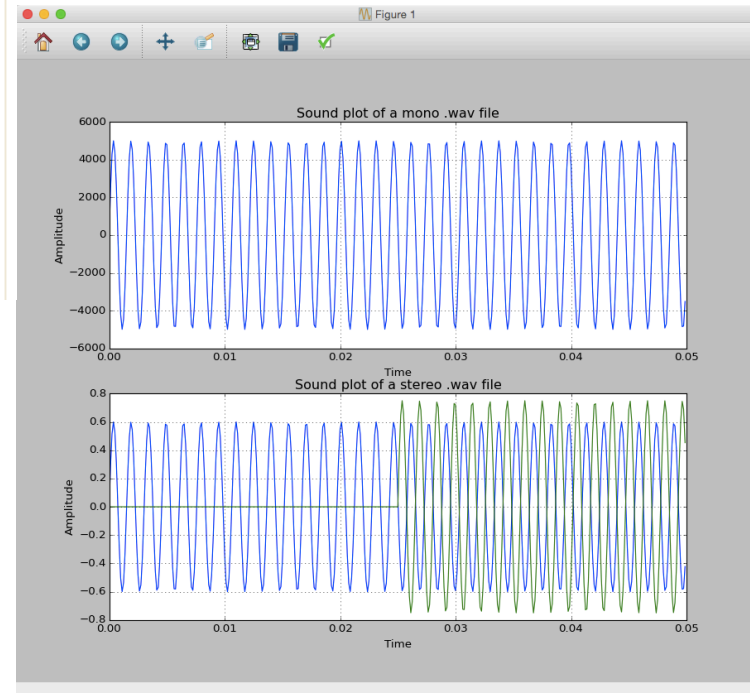
```
130 ## Example 3 - manipulating sounds - creating pseudo-stereo from mono:
131 from numpy import zeros, concatenate
132 from scipy import fft, arange, ifft, sin, pi
133 from scipy.io.wavfile import read, write
134
135 (Fs, x) = read('files/lecture8/melody.wav')
136 x      # contains all the samples
137 y = x  # we create the second channel
138 z=zeros([200]) # create an array of zeros
139 # 1. Time/Phase shift the two channel:
140 L=concatenate((x,z)) # we add zeros after the 'x' signal to create Left channel
141 R=concatenate((z,y)) # we add zeros before the 'y' signal to create Right channel
142 A=zeros([len(L),2]) # we now create the array to store 'x' and 'y' as L and R
143 A[:,0]=L # we assign the Left channel
144 A[:,1]=R # we assign the Right channel
145 # 2. Amplitude change:
146 A[:,0]=A[:,0]*1.2e-4 # we decrease the amplitude on the Left to avoid clipping
147 A[:,1]=A[:,1]*1.5e-4 # ampl. decrease on Right channel is more since it is delayed
148
149 # we write the file:
150 write('files/lecture8/melody_stereo.wav',Fs,A)
```


Working with files

- Working with files – sound 4/4
 - we plot the mono and stereo music we just created

```
152 ## Lets plot the mono and pseudo-stereo sounds:
153 from pylab import linspace, plot, subplot, title, xlabel, ylabel, grid, pause
154
155 length = len(L) # number of samples in either channel 'L' (they are equal)
156 time = length/Fs # calculate the length of the .wav file in seconds
157 t = linspace(0,time,length) # create evenly spaced numbers between [0:time]
158
159 subplot(2,1,1)
160 plot(t[0:400],L[0:400]) # plot the first 400 samples from the mono signal 'L'
161 title('Sound plot of a mono .wav file')
162 xlabel('Time'); ylabel('Amplitude'); grid(True)
163
164 subplot(2,1,2)
165 plot(t[0:400],A[0:400]) # plot the first 400 samples from the stereo signal 'A'
166 title('Sound plot of a stereo .wav file')
167 xlabel('Time'); ylabel('Amplitude'); grid(True)
168 pause(1)
```

manipulating
.wav files



The Fast Fourier Transform

- The Fast Fourier Transform – quick intro
 - FFT is the faster implementation of the DFT
 - FFT is the basis for frequency analysis that converts any signal in time to frequency domain
 - fft is the Fast Fourier Transform (FFT) converts time-domain signals to frequency-domain
 - ifft is the Inverse Fast Fourier Transform (IFFT) and converts frequency to time-domain
 - the Fourier transform is represented like this (ex: an audio signal):

$$X_k = \sum_{n=0}^{N-1} x_n e^{-i2\pi k \frac{n}{N}} \quad \text{or} \quad \hat{f}(\xi) = \int_{-\infty}^{\infty} f(x) e^{-2\pi i x \xi} dx$$

where: $[x_0, \dots, x_{N-1}]$ are complex conjugate numbers and $[k=0, \dots, N-1]$

- the most commonly used FFT is the Cooley–Tukey algorithm

The Fast Fourier Transform

- The Fast Fourier Transform – quick intro
 - FFT is the faster implementation of the DFT
 - FFT is the basis for frequency analysis that converts any signal in time to frequency domain
 - fft is the Fast Fourier Transform (FFT) converts time-domain signals to frequency-domain
 - ifft is the Inverse Fast Fourier Transform (IFFT) and converts frequency to time-domain
 - the Fourier transform is represented like this (ex: an image):

$$F(k, l) = \sum_{i=0}^{N-1} \sum_{j=0}^{N-1} f(i, j) e^{-i2\pi(\frac{ki}{N} + \frac{lj}{N})}$$

where: where $f(a, b)$ is the image in the spatial domain and $F(k, l)$ corresponds to each pixel

- the most commonly used FFT is the Cooley–Tukey algorithm

The Fast Fourier Transform

- The Fast Fourier Transform

Example:

- notice that we only import what we need

- we add three simple tones to create one complex tone

- to go from time domain to frequency domain we use `fft` in two ways

- each frequency bin represents frequency in [Hz] ->

```
243 # FFT example:
244 from numpy.fft import rfft
245 from scipy import arange, sin, pi, fft, real, imag, log10
246 from scipy.io.wavfile import write
247 from matplotlib.pyplot import figure, plot, subplot, axis, grid
248 from pylab import xticks, yticks, xlim, ylim, xlabel, ylabel, title
249
250 Fs=4000 # sampling frequency
251 a = arange(1024) # create a vector holding the number of bins
252 signal1 = sin(2*pi*a*(650/Fs)) # create a tone with frequency = 650Hz
253 signal2 = sin(2*pi*a*(1150/Fs)) # create a tone with frequency = 1.15kHz
254 signal3 = sin(2*pi*a*(1450/Fs)) # create a tone with frequency = 1.425kHz
255 signal4 = sin(2*pi*a*(1250/Fs)) # create a tone with frequency = 1.25kHz
256
257 # Create a complex tone:
258 signal = signal1 + signal2 + signal3
259
260 # Take the FFT of the complex signal:
261 freq_domain_npy = rfft(signal) # using npy
262 freq_domain_cpy = fft(signal) # using cpy
263 bin_val = Fs/len(a) # calculate the value of each frequency bin in [Hz]
264 bin_val # each bin in [Hz]
265 t = arange(1,Fs/2+2,bin_val) # create a vector of frequency bins in [Hz]
266
267 # Save into a wav file:
268 write('files/lecture8/fft_file_example.wav',Fs,signal) # save to file
```



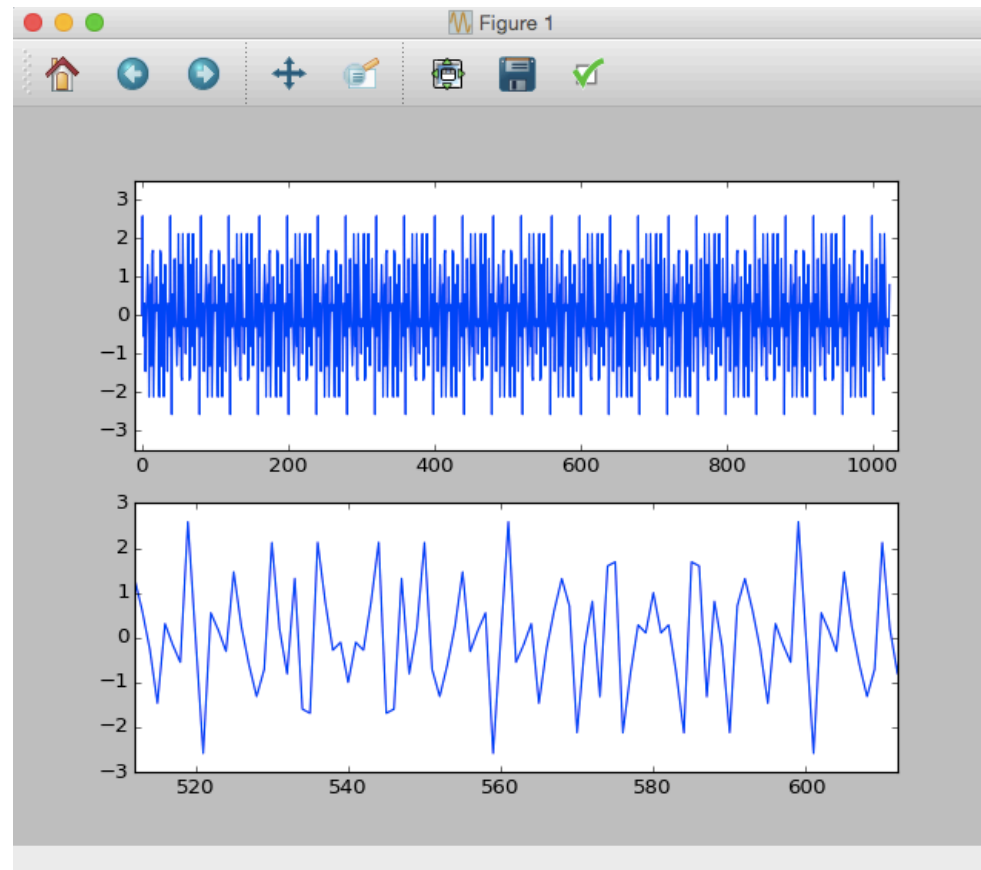
```
In [113]: bin_val # each bin in [Hz]
Out[113]: 3.90625
```

The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
270 # Plot the raw Time domain signal:  
271 figure(1), subplot(2,1,1), plot(signal), xlim(-10, len(a)+10), ylim(-3.5,3.5)  
272 subplot(2,1,2), plot(signal), xlim(len(a)/2, len(a)/2+100) # just a snipped of the signal
```



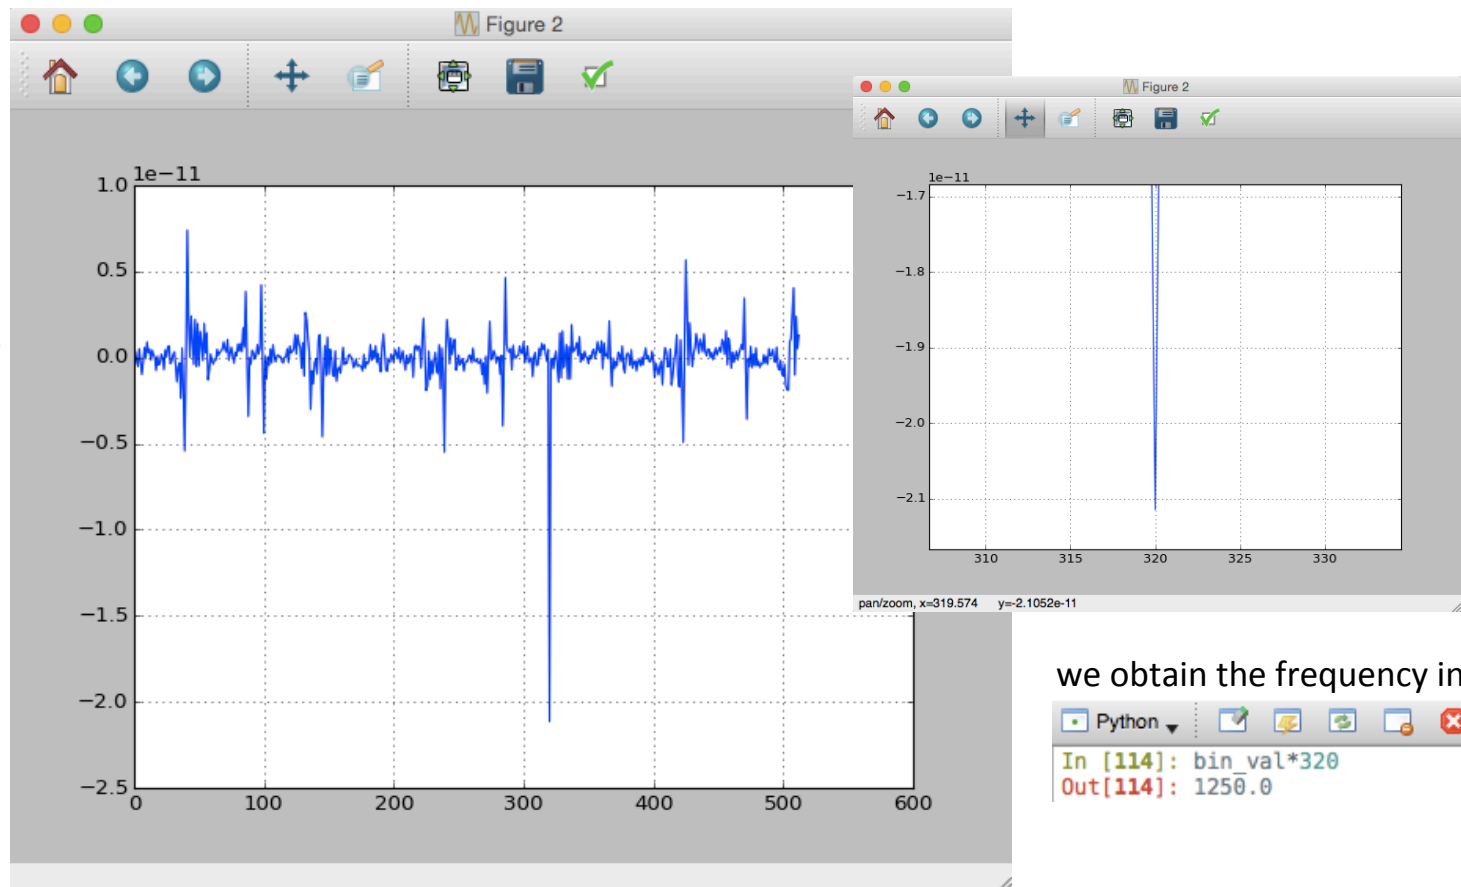
The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
274 # Plot the Frequency domain of 'signal4':  
275 figure(2), plot(rfft(signal4)) # observe the quantization noise due to rounding errors  
276 grid(True)
```

- observe the
quantization noise
due to rounding
errors



we obtain the frequency in [Hz]:

```
Python  
In [114]: bin_val*320  
Out[114]: 1250.0
```

The Fast Fourier Transform

- The Fast Fourier Transform


Example:

```
278 # Plot the Frequency domain signal using numpy:
279 figure(3)
280 subplot(3,1,1),
281 title('Frequency domain real part of the signal - raw frequency bins'),
282 plot(freq_domain_npy),
283 xlim(0,len(a)/2),
284 ylabel('amplitude'),
285 grid(True)
286
287 # Plot the magnitude of the complex signal output:
288 subplot(3,1,2),
289 plot(abs(freq_domain_cpy)),
290 title('Magnitude of the complex signal - raw frequency bins'),
291 xlim(0,len(a)/2),
292 ylabel('amplitude'),
293 grid(True)
294
295 # Plot in dB scale:
296 subplot(3,1,3),
297 plot(t,10*log10(freq_domain_npy)),
298 xlabel('frequency in Hz'),
299 ylabel('amplitude in dB'),
300 axis('tight'),
301 grid(True)
```

- notice the difference
in the x scales

```
In [11]: bin_val*166
Out[11]: 648.4375
In [12]: bin_val*294
Out[12]: 1148.4375
In [13]: bin_val*371
Out[13]: 1449.21875
```

create a vector
*(650/Fs))
*(1150/Fs))
*(1450/Fs))
*(1250/Fs))



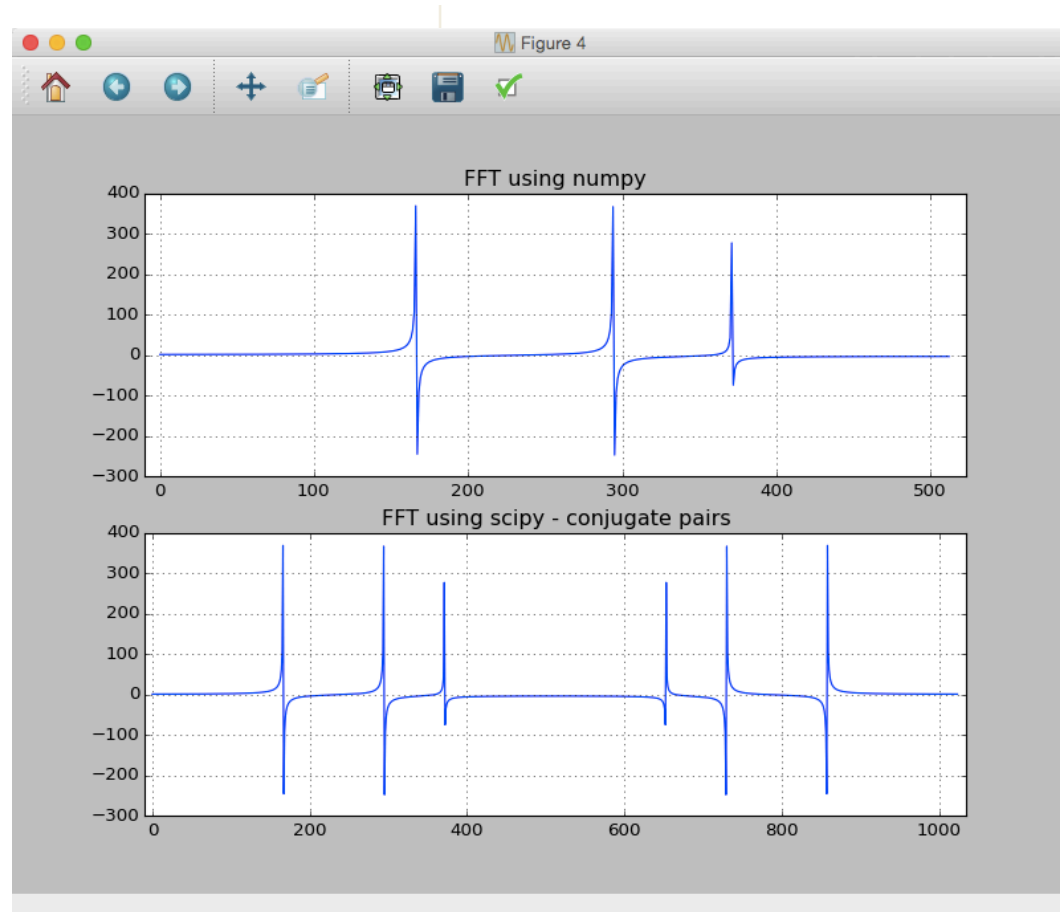
The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
302 # Plot the Frequency domain signal using numpy:
303 figure(4)
304 subplot(2,1,1),
305 plot(freq_domain_npy),
306 xlim(-10, len(t)+10),
307 title('FFT using numpy'),
308 grid(True)
309
310 # Plot the Frequency domain signal using scipy:
311 subplot(2,1,2),
312 plot(freq_domain_cpy),
313 xlim(-10, len(freq_domain_cpy)+10),
314 title('FFT using scipy - conjugate pairs'),
315 grid(True)
384 pause(1)
```

- notice the difference between:
 - `rfft` from NumPy and
 - `fft` from Scipy



Signal Processing

- Signal Processing – sound processing: spectrogram
 - spectrogram is a 3-D way of visualizing the frequency domain of any given signal
 - spectrograms are 3-D because they represent frequencies and their magnitudes over time in a given signal
 - signals are usually: sounds, music and speech, but can also be image signals
 - sometimes spectrograms are referred to as waterfalls, voiceprints, or voicegrams
 - they are the perfect tool for phonetic analysis in visualizing spoken words
 - spectrogram visualizing functionality can be uploaded in one of two ways:
 - In [1]: `from pylab import specgram` ... or
 - In [1]: `from matplotlib.pyplot import specgram`
 - sometimes the spectrogram format varies and the vertical and horizontal axes can be switched
 - spectrograms are usually generated in two ways, by using:
 - FFT calculated from a given time signal
 - Filterbanks resulting from a sequence of bandpass filters

Signal Processing

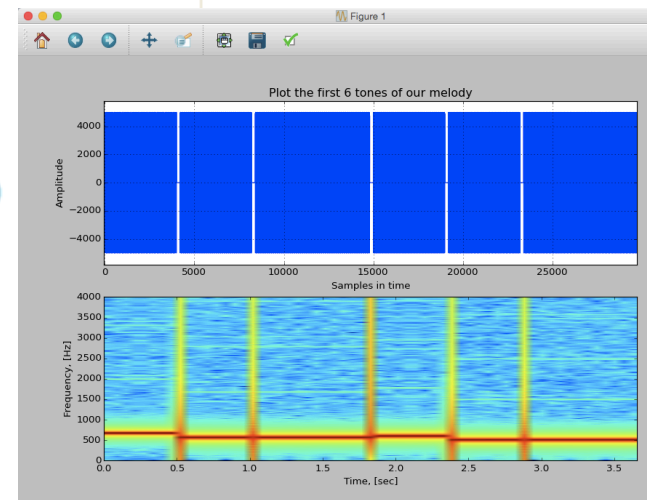
- Signal Processing – sound processing: spectrogram
 - here are some of the parameters users have control over:
 - F_s – the sampling frequency calculating the Fourier frequencies, in cycles per time
 - $NFFT$ – the amount of frequency bins represented in each FFT window
 - window type used – Hamming, Hanning, Bratlett, Blackman, Kaiser
 - noverlap – amount of overlap between window blocks. the default overlap is 128 samples
 - mode – type of spectrogram to visualize: { 'psd' | 'magnitude' | 'angle' | 'phase' }, where:
 - psd – is the power spectral density
 - magnitude – is the magnitude spectrum
 - angle – represents the phase spectrum without unwrapping
 - phase - is the phase spectrum with unwrapping
 - scale – how the data should be displayed { 'default' | 'linear' | 'dB' }, where:
 - default – linear
 - linear – means no scaling will be used
 - dB – when mode=psd' the dB scale is $(10 \cdot \log_{10})$, otherwise it is $(20 \cdot \log_{10})$
 - F_c – the center frequency can be controlled
 - cmap – is the colormap chosen

Signal Processing

- Signal Processing – sound processing: spectrogram

Example:

```
3 # Signal Processing:
4 ## 1. Spectrogram:
5
6 from pylab import specgram, plot, subplot, title, xlabel, ylabel, grid, axis, xlim, ylim, pause
7 from scipy.io.wavfile import read
8
9 (Fs, x) = read('files/lecture9/melody.wav') # Fs - sampling frequency, x - signal
10
11 # Lets plot signal 'x':
12 subplot(2,1,1)
13 plot(x)
14 xlim([-50,29750]); ylim([-5800,5800]) # limit it to the first 6 tones only
15 title('Plot the first 6 tones of our melody')
16 xlabel('Samples in time'); ylabel('Amplitude')
17 grid(True)
18
19 # Now we plot the spectrogram of the first 6 tones:
20 subplot(2,1,2)
21 xlabel('Time, [sec]'); ylabel('Frequency, [Hz]')
22 specgram(x[0:29750], NFFT=512, Fs=Fs, noverlap=64, mode='magnitude')
23 pause(1)
```



Signal Processing

- Signal Processing – sound processing: spectrogram

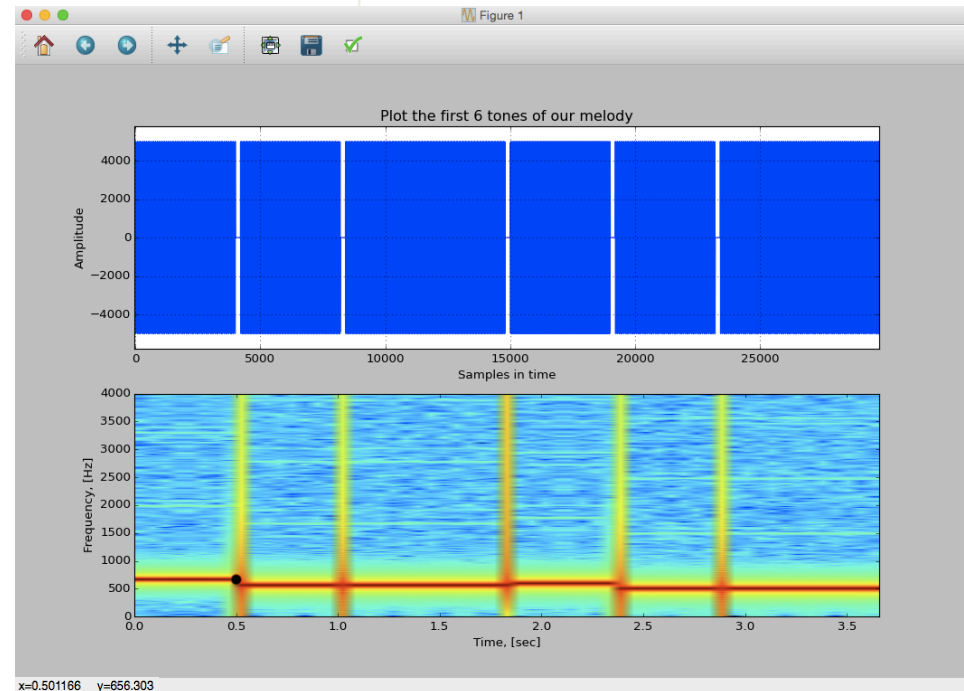
Example:

lets recall the code snipped from the melody:

```
97 # Creating each tone with Fs samples per second and length 0.5 or 0.8 seconds:
98 Es = note(659,Fs,0.5,amplitude=5000)
99 El = note(659,Fs,0.8,amplitude=5000)
100 C#s = note(554,Fs,0.5,amplitude=5000)
101 C#l = note(554,Fs,0.8,amplitude=5000)
102 D = note(587,Fs,0.5,amplitude=5000)
103 Bs = note(494,Fs,0.5,amplitude=5000)
104 Bl = note(494,Fs,0.8,amplitude=5000)
105 As = note(440,Fs,0.5,amplitude=5000)
106 Al = note(440,Fs,0.8,amplitude=5000)
```

and the melody: E C# C# D B B A B C# D E E E
E C# C# D B B A C# E E A A A

- notice how the frequency tones are represented by the 'hottest' red line in the spectrum
- we also notice how the timing of each tone corresponds to 0.5 and 0.8 sec
- for E: x=0.501166 [sec], y=656.303 [Hz]



Signal Processing

- Signal Processing – sound processing: speech

Vowels and phonemes
in American – English

We use five letters to
represent the vowel sounds:
a, e, i, o, u

Words [↗]	Ladefoged (2006) [↗]	Roach (2009) [↗]	Words [↗]	Ladefoged (2006) [↗]	Roach (2009) [↗]
<u>f</u> ee★ [↗]	/i/ [↗]	/i:/ [↗]	b <u>i</u> rd★ [↗]	/ɜ/, /ɜː/ [↗]	/ɜ:/ [↗]
<u>h</u> ard★ [↗]	/ɑ/ [↗]	/ɑː/ [↗]	b <u>e</u> d★ [↗]	/ɛ/ [↗]	/e/ [↗]
<u>f</u> ood★ [↗]	/u/ [↗]	/u:/ [↗]	<u>a</u> ttend [↗]	/ə/ [↗]	/ə/ [↗]
<u>l</u> ord★ [↗]	/ɔ/ [↗]	/ɔ:/ [↗]	<u>b</u> ook [↗]	/u/ [↗]	/u/ [↗]
<u>h</u> ot★ [↗]	/ɑ/(GA), [↗] /ɒ/ (RP) [↗]	/ɒ/ [↗]	<u>g</u> o★ [↗]	/ou/(GA) [↗] /əu/(RP), [↗]	/əu/ [↗]
<u>b</u> us, [↗]	/ʌ/ [↗]	/ʌ/ [↗]	<u>b</u> oy [↗]	/ɔɪ/ [↗]	/ɔɪ/ [↗]
<u>b</u> ook [↗]	/u/ [↗]	/u/ [↗]	<u>b</u> ig [↗]	/ɪ/ [↗]	/ɪ/ [↗]
<u>d</u> ear, [↗]	/ɪə/ [↗]	/ɪə/ [↗]	<u>c</u> are★ [↗]	/ɛə/ [↗]	/eə/ [↗]
<u>b</u> ike [↗]	/aɪ/ [↗]	/aɪ/ [↗]	<u>h</u> ow [↗]	/aʊ/ [↗]	/aʊ/ [↗]
<u>c</u> ake [↗]	/eɪ/ [↗]	/eɪ/ [↗]	<u>t</u> our [↗]	/ʊə/ [↗]	/ʊə/ [↗]
★ =There is a difference between two systems [↗]					

IPA fonts in Ladefoged (2006) and Roach (2009)

Signal Processing

- Signal Processing – sound processing: speech


Formants are resonances in the vocal tract

Each vowel is formed by formants: a concentration of acoustic energy around a specific frequency

Each formant has a different center frequency with higher amplitude

Formants for different genders or kids vary

Spectrograms of the American English Vowels



(Ladeforged 2006:185-187)

Signal Processing

- Signal Processing – sound processing: speech


Spectrograms of the American English Vowels

Formants are resonances in the vocal tract

Each vowel is formed by formants: a concentration of acoustic energy around a specific frequency

Each formant has a different center frequency with higher amplitude

Formants for different genders or kids vary



(Ladefoged & Johnson, 2011:193)

Signal Processing

- Signal Processing – sound processing: speech


Formants are resonances in the vocal tract

Each vowel is formed by formants: a concentration of acoustic energy around a specific frequency

Each formant has a different center frequency with higher amplitude

Formants for different genders or kids vary

Spectrograms of the American English Vowels



(Ladefoged & Johnson, 2011:193)

Signal Processing


- Signal Processing – sound processing: speech

Formant frequencies for each vowel

We use five letters to
represent the vowel sounds:
a, e, i, o, u

use LPC for formant
estimation from **Audiolazy**:

`pip install audiolazy`



Vowel	F1(Hz)	F2(Hz)	F3(Hz)
i:	280	2620	3380
ɪ	360	2220	2960
e	600	2060	2840
æ	800	1760	2500
ʌ	760	1320	2500
ɑ:	740	1180	2640
ɒ	560	920	2560
ɔ:	480	760	2620
ʊ	380	940	2300
u:	320	920	2200
ɜ:	560	1480	2520

Adult male formant frequencies in Hertz collected by J.C.Wells around 1960.

Note how F1 and F2 vary more than F3.

Signal Processing

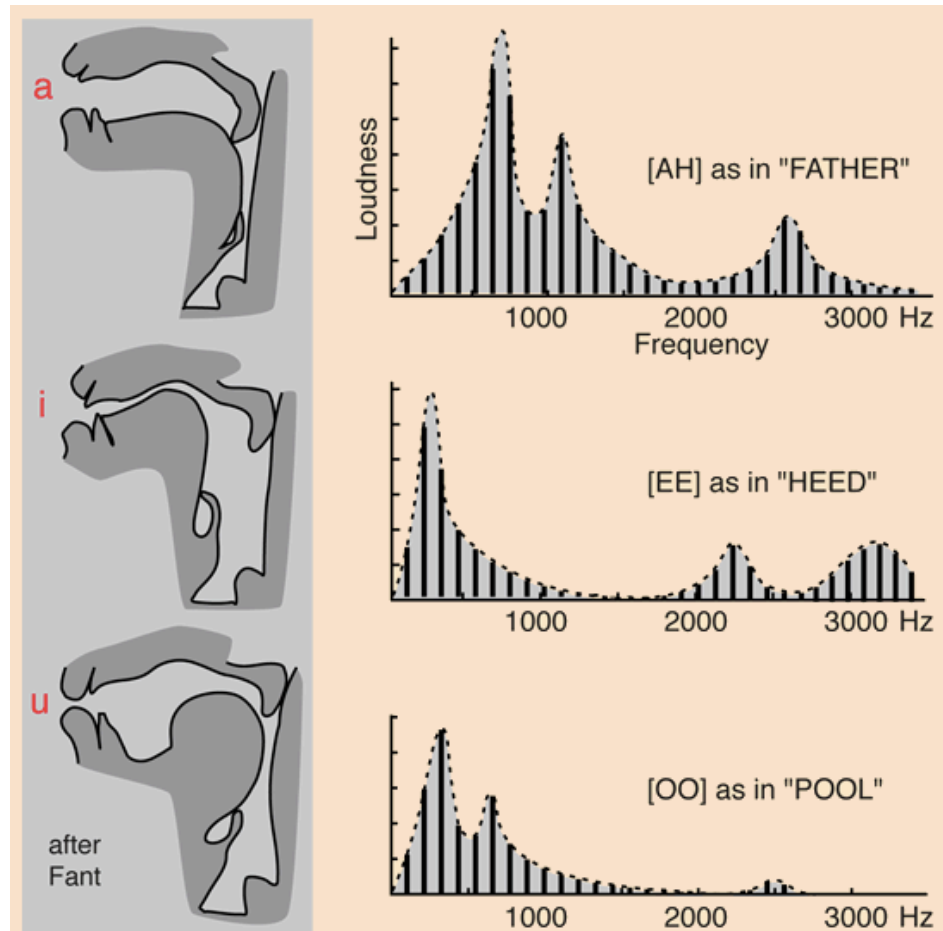
- Signal Processing – sound processing: speech

See levels for each formant

Each vowel is formed by formants

Each formant has a different frequency and amplitude

Formants for different genders or kids vary



Feature Extraction – glottal signal

The Glottal Signal:

conveys speaker identity, mode of speaking

airflow passing through the glottis creates voiced sounds: *vowels, semivowels, nasals, diphthongs, consonants*




Glottal waveform is greatly affected by the emotional state -
Laukkanen et al., 1996

MRI - collected from **Centre for Speech Technology**
and l'Institute da la Communication Parlee in Grenoble

Feature Extraction – glottal signal

Speech production model:

3 concatenated linear time-varying subsystems



$$S(z) = G(z)V(z)R(z)$$

$$G(z) = \frac{S(z)}{V(z)R(z)}$$

